

Homework 3: Solutions Dominik Gresch

Exercise 1: Cache mechanics

a) The general solution is given by

$$M = \frac{1}{2(bs+t)} \left(s + \lceil t/b \rceil + \sum_{j=0}^{\lceil t/b \rceil - 1} \Theta \left(\sum_{i=(j+1)b-1}^{bs+t-1} \delta_{\beta_i, s+j} \right) \right. \\ \left. + \sum_{j=0}^{\lceil t/b \rceil - 1} \Theta \left(\sum_{i=\min[b(s+j+1)-1, bs+t-1]}^{2bs+t-1} \delta_{\beta_i, j} \right) + \sum_{j=0}^{\lceil t/b \rceil - 1} \delta_{\alpha_i, \beta_{i-1}} \cdot (1 - \delta_{\gamma_i, 0}) \right)$$

where

$$\Theta(x) = \begin{cases} 0 & \text{if } x \leq 0 \\ 1 & \text{else} \end{cases}$$

$$\alpha_i = \lfloor (bs + ((bi) \% (bs+t))) / b \rfloor \% 2s$$

$$\beta_i = \lfloor ((bi) \% (bs+t)) / b \rfloor \% 2s$$

$$\gamma_i = (i \% (bs+t)) \% 2s$$

and δ is the Kronecker Delta. α_i and β_i describe in which set y and x are at iteration i , δ_i the block offset of y .

The first 2 terms come from the misses that happen because y gets loaded into cache the first time.

The third term comes from cacheblocks at the beginning of y that have been evicted between the time when y first accessed them and the time y accesses them again (second part of the loop).

The fourth term is the same, but with the end of y that overlaps with x (in terms of cache address).

Finally, the last term describes cache misses due to a “race condition”, between x and y , those are misses where x and y are at the same cache location at the same time (but not when y accesses the first element of a block, because those are already accounted for).

I have provided a **Mathematica** file for evaluation: see `/ex01/ex01.nb`

As a last remark, for (s, b, t) which fulfil the condition $t \leq b \leq s$ or $b + t \leq s$ (I only tested these limits numerically), this term becomes much easier:

$$M = \frac{s + \lceil t/b \rceil}{2(bs+t)}$$

b) x : for the first 7 i , there are only misses (access the first element of a cache block each time). After that, there will be 6 hits (accessing the second element). For $i \in \{13, 25\}$, there will be misses only for the elements that have been evicted due to conflicts with y (i.e. $i = 13, 19$)

$\Rightarrow x$: MMMMMMMHHHHHH MHHHHMHHHHHH

y: For the first 13 i, there are alternating misses and hits. After that, there will be a miss only for i = 13 and i = 25 (the elements that have been evicted due to conflict with x).

$\Rightarrow y : \text{MHMHMHMHMHMHMH MHMMMMMMMMMMMM}$

c) Since `sum` is in a register, the traffic will be exactly $2 \text{ doubles} \cdot \# \text{misses}$. There are 18 misses, hence

$$Q = 18 \cdot 2 \cdot 8 \text{ bytes} = 288 \text{ bytes}$$

There are 2 flops in each loop iteration, 26 iteration steps

$$\Rightarrow W = 52 \text{ flops} \Rightarrow I = \frac{13}{72} \frac{\text{flop}}{\text{byte}}$$

Exercise 2: Cache mechanics

- a) The upper bound for I is determined by the minimum number of bytes read (reading the three variables exactly once) and the number of flops.

$$\begin{aligned} Q &\geq (12 + 4 + 3) \text{ doubles} = 152 \text{ bytes} \\ W &= 12 \text{ add} + 12 \text{ mult} = 24 \text{ flop} \\ \Rightarrow I &\leq \frac{12 \text{ flop}}{76 \text{ byte}} \end{aligned}$$

- b) In every iteration of the innermost loop, there is first an access to $A[i][j]$ and $x[j]$ and then an access to $y[i]$ (the multiplication gets evaluated first).

Each cache block fits exactly one **double**.

Looking only at A and x , each even j fills up the 0 - th set and each odd j fills up the 1st set. Since A does not use the same value twice and x only uses the same value after 4 j (when all previous values have already been evicted), there are no cache hits for either A or x . This results in 12 misses for A and 12 misses for x .

For y , there are two different scenarios. If $i \% 2 == 0$ (i.e. $y[i]$ is in the first set), the hit/miss pattern for $j = 0, \dots, 3$ is MHMH (2 misses). If $i \% 2 == 1$ (i.e. $y[i]$ is in the second set), the miss/hit pattern is MMHM (3 misses). Because the first scenario happens twice and the second once, this gives 7 misses for y .

$$M_{L_1} = \frac{12 + 12 + 7 \text{ misses}}{36} = \frac{31 \text{ misses}}{36 \text{ access}}$$

- c) Every read miss creates 8 **bytes** of traffic. Additionally, y has to be written to memory exactly the same number of times as misses for y (each time the value gets brought in it will eventually have to be written back).

$$Q = (31 + 7) \cdot 8 \text{ bytes} = 304 \text{ bytes}$$

of traffic, and hence

$$I_{L_1} = \frac{3 \text{ flop}}{38 \text{ byte}}$$

- d) The variables all fit into the L_2 - cache (no conflicts): A takes up one block per set for the first 6 sets, x and y fit into the two blocks of the 7-th and 8-th set respectively.

The only misses come from bringing the variables in once. Since each block contains 2 **doubles**, this gives rise to

- $\frac{12}{2} \text{ misses} = 6 \text{ misses}$ for A
- $\frac{4}{2} \text{ misses} = 2 \text{ misses}$ for x
- $\lceil \frac{3}{2} \rceil \text{ misses} = 2 \text{ misses}$ for y

$$\Rightarrow 10 \text{ misses} \Rightarrow M_{L_2} = \frac{10 \text{ misses}}{3 \cdot 12 \text{ access}} = \frac{5 \text{ misses}}{18 \text{ access}}$$

- e) Memory traffic comes from reading each variable and writing y . For y , there is an additional **double** of memory traffic since it doesn't fit perfectly into the second block (effectively, 4 **double** of traffic for each complete read / write of y).

$$\Rightarrow Q = (12 + 4 + 4 \cdot 2) \cdot 8 \text{ byte} = 192 \text{ byte} \Rightarrow I_{L_2} = \frac{1 \text{ flop}}{8 \text{ byte}}$$

Exercise 3: Associativity

I do not agree with your statement. Counterexample:

Let x , y and z be 3 arrays of size $B/8$ (they fit exactly into one block), and their addresses such that (for the direct mapped cache), x gets mapped into the 0 - th set, y and z into the $S/2$ - th. Of course, for the 2 - way associative cache, they will all be mapped into the 0-th cache set. Now consider the following code (assuming `sum` is in registers, otherwise cold cache):

```
double sum = 0;
for(int i = 0; i < B / 8; ++i) {
    sum += x[i];
    sum += y[i];
    sum += z[i];
}
```

Whilst the direct mapped cache produces hits for x (after the initial miss), and misses only for y and z , the 2 - way associative cache produces only misses.

direct mapped: MMMHMMHMM...

2 - way associative: MMMMMMMMM...

The reason for this is that x stays in cache in the direct mapped case whilst it gets evicted (because it is least recently used) in the 2 - way associative case.




Figure 1: first loop iteration with direct mapped cache




Figure 2: first loop iteration with 2 - way associative cache